

DS_A1_provenance_audit

September 30, 2026

1 DS-A1 · Data Provenance and Measurement Audit

1.0.1 Dataset: lerobot/svla_so101_pickplace (Hugging Face Hub)

Course: Data Science (UCAS) — Assessed Task DS-A1

Student: Jin Xinhui

Date: 2026/9/27 final run 2026-09-30

Repository: <https://github.com/lynn3214/ds-a1-data-provenance-audit>

Submitted version: tag v1.0-submission

1.1 1. Reproducibility setup

We fix the random seed, record package versions, and record the Python version and OS platform the notebook was run on. The printed output below is also recorded in `data_card/SOURCE_LICENSE_HASH.md` under “Execution environment”, and the full `pip freeze` is written to `environment_freeze.txt` by the next cell.

```
[1]: import sys, os, platform, random, json, hashlib, subprocess
    from pathlib import Path

    SEED = 42
    random.seed(SEED)

    import numpy as np
    np.random.seed(SEED)

    # Third-party packages used throughout this notebook.
    # If any import fails, run:
    # pip install -r ../requirements.txt
    import pandas as pd
    import pyarrow.parquet as pq
    import matplotlib.pyplot as plt
    import matplotlib.patches as mpatches
    from sklearn.neighbors import KNeighborsRegressor
    from sklearn.model_selection import GroupShuffleSplit, ShuffleSplit
    from sklearn.metrics import mean_squared_error
```

```

from huggingface_hub import HfApi, snapshot_download, hf_hub_download

print("Python      :", sys.version.replace("\n", " "))
print("Platform    :", platform.platform())
print("pandas       :", pd.__version__)
print("numpy         :", np.__version__)
import pyarrow, sklearn, matplotlib, huggingface_hub
print("pyarrow       :", pyarrow.__version__)
print("scikit-learn   :", sklearn.__version__)
print("matplotlib     :", matplotlib.__version__)
print("huggingface_hub:", huggingface_hub.__version__)
print("SEED          :", SEED)

```

```

Python      : 3.10.18 | packaged by conda-forge | (main, Jun  4 2025, 14:46:00)
[Clang 18.1.8 ]
Platform    : macOS-15.5-arm64-arm-64bit
pandas      : 2.3.1
numpy       : 2.2.6
pyarrow     : 25.0.1
scikit-learn : 1.7.1
matplotlib  : 3.10.5
huggingface_hub: 2.0.0
SEED        : 42

```

```

[2]: # Freeze the full dependency list actually installed in this environment.
     # The resulting environment_freeze.txt is committed as part of the
     # reproducibility record.
     REPO_ROOT = Path("..").resolve()
     freeze_path = REPO_ROOT / "environment_freeze.txt"
     with open(freeze_path, "w") as f:
         subprocess.run([sys.executable, "-m", "pip", "freeze"], stdout=f)
     print(f"Wrote {freeze_path}")

```

Wrote /Users/jinshenhui/Documents/UCAS/DS/assignments/ds-a1-lerobot-pickplace-audit/environment_freeze.txt

Execution order note: this notebook has a strict top-to-bottom variable dependency chain (e.g. Section 7 defines `action_arr/state_arr/JOINT_NAMES` used by Sections 8, 11, 12; Section 9 defines `df_sorted` used by Sections 10, 11, 12). Running cells out of order, or only running a subset, will raise `NameError`. All randomness is seeded (`SEED=42`, Section 1; `random_state=SEED` in Section 11’s splitters), so a full, in-order rerun is expected to reproduce identical numeric results every time.

1.2 2. Question definition

Stakeholder. An embodied-AI / robot-learning researcher deciding whether `lerobot/svla_so101_pickplace` can be used as a benchmark or fine-tuning dataset for an imitation-learning (behavior cloning) policy, and how to split it for evaluation.

Testable question. *Does the recording structure of this dataset (one robot type, one task string, 50 episodes; environment and operator variation are undisclosed) support a **frame-level** random train/test split for evaluating an action-prediction baseline, or does it require an **episode-level** (grouped) split to avoid leakage? How large is the resulting optimism bias if the wrong split is used?*

Population. All possible demonstration trajectories of an SO-100/SO-101 follower arm performing the single task recorded in `meta/tasks.parquet` (“pink lego brick into the transparent box”) under broadly similar lab conditions.

Sample. The 50 recorded episodes (11,939 frames total) published in this dataset snapshot.

Unit of analysis. A single timestep / frame (30 fps). The secondary/grouping unit is the episode (`episode_index`, 0–49).

Target / response variable. `action`: the 6-dimensional commanded joint position vector (`shoulder_pan.pos`, `shoulder_lift.pos`, `elbow_flex.pos`, `wrist_flex.pos`, `wrist_roll.pos`, `gripper.pos`).

Estimand. The expected held-out MSE (in the dataset’s action units, squared; the unit is undocumented) of a k-NN regressor mapping `observation.state` $\rightarrow$ `action`, when the model is applied to a **new episode** that was not used in training. Two estimators of this quantity are compared: - $\hat{\theta}_{frame}$: MSE under a random frame-level 80/20 split (frames from the same episode can appear in both train and test). - $\hat{\theta}_{episode}$: MSE under a grouped 80/20 split by `episode_index` (an entire episode is either fully in train or fully in test).

Because the target is performance on a new episode, $\hat{\theta}_{episode}$ is the estimator that matches the estimand, and $\hat{\theta}_{frame}$ is expected to be biased low if frames within an episode are strongly dependent. The quantity reported is the gap $\Delta = \hat{\theta}_{episode} - \hat{\theta}_{frame}$ and the ratio $\hat{\theta}_{episode}/\hat{\theta}_{frame}$. A large positive Δ is *consistent with* the frame-level split being optimistically biased. That this is *caused by* near-duplicate temporal neighbors is a hypothesis; it is not tested directly in this notebook (see Section 15).

Explicit boundaries. This audit does **not** attempt to evaluate whether a trained policy would succeed physically (no success/reward label exists in this dataset — see Section 8). It only evaluates whether the *evaluation protocol itself* is trustworthy.

1.3 3. Data-generating process (DGP) and provenance lineage

The diagram below is drawn as an actual matplotlib figure (not a Mermaid block) so it renders identically in the notebook, the exported PDF, and on GitHub.

Declared facts, as read from the pinned snapshot’s own `meta/info.json` (not from the rendered Hub webpage; verified in Section 4-5): `robot_type = so100_follower`, `fps = 30`, `total_episodes = 50`, `total_frames = 11939`, `total_tasks = 1`, `codebase_version = v3.0`, `license = Apache-2.0`.

Provenance/version finding: before the snapshot was inspected, my planning-stage assumption was a v2.1-style layout (a single `meta/episodes.jsonl` file, one parquet file per episode). The pinned snapshot (f641879e22172be7e8161d5e6c1503c2d2feb657), however, declares `codebase_version = v3.0` in its own `meta/info.json` and uses a different on-disk layout: `meta/episodes/` as sharded parquet, `meta/tasks.parquet`, multi-episode parquet chunks under `data/`, and multi-episode video files (Section 5 shows several episodes sharing one video file via

from_timestamp/to_timestamp). At freeze time the Hub HEAD was identical to the pinned commit (Section 4), so this disagreement is not explained by the page showing a newer commit; I did not establish where the v2.1 assumption originated. The pinned revision’s own `meta/info.json` is therefore used as the ground truth for this audit’s schema and format claims. This illustrates why DS-A1 asks for freezing and hashing the actual downloaded snapshot instead of relying on a description of it.

Provenance gaps (explicitly not resolved by the dataset card — do not silently assume an answer): - Number of distinct human teleoperators is undisclosed. If a single operator recorded all 50 episodes, the action distribution reflects one person’s motor style, not a population of operators. - Whether object placement / lighting / camera position varied across episodes, or was held approximately fixed, is undisclosed in the card. - No ground-truth “task success” or “grasp success” label is included in the schema — only the commanded trajectory is recorded, not the outcome.

These gaps are treated as findings, not filled in with assumptions.

```
[3]: fig, ax = plt.subplots(figsize=(10, 3.6))
ax.axis("off")

stages = [
    "Human teleoperator\n(count undisclosed)\ndrives leader arm\n(inferred)",
    "so100_follower arm\nexecutes; joint state\nand action are logged",
    "2 cameras (up, side)\nvideo @ 30 fps\n(declared in info.json)",
    "LeRobot format v3.0\nparquet (per-frame) +\nmp4 (multi-episode files)",
    "Hub: lerobot/\nsvla_so101_pickplace\n(pinned snapshot)",
]

n = len(stages)
box_w, box_h = 1.7, 1.6
gap = 0.55
x0 = 0.2
y0 = 0.3

for i, text in enumerate(stages):
    x = x0 + i * (box_w + gap)
    rect = mpatches.FancyBboxPatch((x, y0), box_w, box_h,
                                   boxstyle="round,pad=0.06",
                                   linewidth=1.2, edgecolor="#333333",
                                   facecolor="#eaf2fb")

    ax.add_patch(rect)
    ax.text(x + box_w / 2, y0 + box_h / 2, text, ha="center", va="center",
           ↪fontsize=8.3)
    if i < n - 1:
        ax.annotate("", xy=(x + box_w + gap - 0.05, y0 + box_h / 2),
                    xytext=(x + box_w + 0.05, y0 + box_h / 2),
                    arrowprops=dict(arrowstyle="->", lw=1.2))


ax.set_xlim(0, x0 + n * (box_w + gap))
```

```

ax.set_ylim(0, y0 + box_h + 0.4)
ax.set_title("Data-generating process: from teleoperation to this notebook",
             ↪fontsize=11)
plt.tight_layout()
fig.savefig("../report/dgp_diagram.png", dpi=180)
plt.show()

```

Data-generating process: from teleoperation to this notebook



1.4 4. Data acquisition and version freezing

We resolve the **exact commit hash** of the dataset repo at the time of first access, then re-download using that *pinned* revision on every subsequent run, so the analysis never silently drifts to a newer version of the dataset.

The commit hash below was resolved on 2026-09-27 and is recorded in `data_card/SOURCE_LICENSE_HASH.md`. Every run, including the submitted one, uses that pinned hash (PINNED_REVISION below); the HEAD lookup in Step 1 is informational only.

```

[4]: REPO_ID = "lerobot/svla_so101_pickplace"
api = HfApi()

# Step 1: look up the current HEAD, purely as a point-in-time record of when
↪this
# audit first observed the repo. This value is NOT used to gate anything below
↪--
# it will drift whenever upstream pushes a new commit, and that drift is
↪expected
# and harmless, not an error.
head_info = api.dataset_info(REPO_ID)
print("Current HEAD commit SHA (informational only, will drift over time):",
      ↪head_info.sha)
print("Current HEAD last modified:", head_info.lastModified)

```

```

# Step 2: the actual frozen version for this audit. Hardcoded once, on first
↳access,
# and never derived from `head_info` again -- this is the real "freeze".
PINNED_REVISION = "f641879e22172be7e8161d5e6c1503c2d2feb657" # frozen on
↳2026-09-27

# Step 3: verify the pinned revision itself is a valid, independently resolvable
# commit. This is the actual reproducibility guarantee: it must succeed today,
↳next
# month, or after upstream has moved far past `head_info.sha` above -- because
↳it
# never compares against head_info at all.
pinned_info = api.dataset_info(REPO_ID, revision=PINNED_REVISION)
print("\nPinned revision resolved successfully:", pinned_info.sha)
print("License tag(s) at pinned revision:",
      pinned_info.card_data.get("license") if pinned_info.card_data else None)
assert pinned_info.sha == PINNED_REVISION, (
    "Pinned revision failed to resolve to itself -- check PINNED_REVISION for
↳typos."
)
print("\nUsing pinned revision for all downstream steps:", PINNED_REVISION)

```

Current HEAD commit SHA (informational only, will drift over time):

f641879e22172be7e8161d5e6c1503c2d2feb657

Current HEAD last modified: 2025-09-27 11:25:41+00:00

Pinned revision resolved successfully: f641879e22172be7e8161d5e6c1503c2d2feb657

License tag(s) at pinned revision: apache-2.0

Using pinned revision for all downstream steps:

f641879e22172be7e8161d5e6c1503c2d2feb657

```

[5]: DATA_DIR = Path("../data_cache")
DATA_DIR.mkdir(exist_ok=True)

local_dir = snapshot_download(
    repo_id=REPO_ID,
    repo_type="dataset",
    revision=PINNED_REVISION,
    local_dir=DATA_DIR / "svla_so101_pickplace",
)
print("Downloaded to:", local_dir)

local_dir = Path(local_dir)
for p in sorted(local_dir.rglob("*"))[:20]:
    print(p.relative_to(local_dir))
print("... (truncated)")

```

```

Downloading bytes:          | 0.00B
Reconstructing (incomplete total...): | 0.00B / 0.00B
Fetching 9 files: 0%|          | 0/9 [00:00<?, ?it/s]

Downloaded to: /Users/jinshenhui/Documents/UCAS/DS/assignments/ds-a1-lerobot-
pickplace-audit/data_cache/svla_so101_pickplace
.cache
.cache/huggingface
.cache/huggingface/.gitignore
.cache/huggingface/CACHEDIR.TAG
.cache/huggingface/download
.cache/huggingface/download/.gitattributes.lock
.cache/huggingface/download/.gitattributes.metadata
.cache/huggingface/download/README.md.lock
.cache/huggingface/download/README.md.metadata
.cache/huggingface/download/data
.cache/huggingface/download/data/chunk-000
.cache/huggingface/download/data/chunk-000/file-000.parquet.lock
.cache/huggingface/download/data/chunk-000/file-000.parquet.metadata
.cache/huggingface/download/meta
.cache/huggingface/download/meta/episodes
.cache/huggingface/download/meta/episodes/chunk-000
.cache/huggingface/download/meta/episodes/chunk-000/file-000.parquet.lock
.cache/huggingface/download/meta/episodes/chunk-000/file-000.parquet.metadata
.cache/huggingface/download/meta/info.json.lock
.cache/huggingface/download/meta/info.json.metadata
... (truncated)

```

```

[6]: for p in sorted((local_dir / "meta").iterdir()):
    kind = "DIR " if p.is_dir() else "FILE"
    print(kind, p.relative_to(local_dir))
    if p.is_dir():
        for sub in sorted(p.rglob("*"))[:5]:
            print("    ", sub.relative_to(local_dir))

```

```

DIR  meta/episodes
      meta/episodes/chunk-000
      meta/episodes/chunk-000/file-000.parquet
FILE meta/info.json
FILE meta/stats.json
FILE meta/tasks.parquet

```

```

[7]: def sha256_of_file(path, chunk_size=1 << 20):
    h = hashlib.sha256()
    with open(path, "rb") as f:
        for chunk in iter(lambda: f.read(chunk_size), b''):
            h.update(chunk)

```

```

    return h.hexdigest()

# meta/          info.json      episodes/tasks/stats
#   chunk    parquet    rglob + is_file()
meta_files = sorted(p for p in (local_dir / "meta").rglob("*") if p.is_file())
print(f"Found {len(meta_files)} files under meta/ (recursively):")
for p in meta_files:
    print(" ", p.relative_to(local_dir))

files_to_hash = meta_files + sorted((local_dir / "data").rglob("*.parquet"))[:2]

hash_records = []
for fp in files_to_hash:
    hash_records.append({
        "relative_path": str(fp.relative_to(local_dir)),
        "size_bytes": fp.stat().st_size,
        "sha256": sha256_of_file(fp),
    })

hash_df = pd.DataFrame(hash_records)
hash_df.to_csv("../data_card/file_hashes.csv", index=False)
hash_df

```

Found 4 files under meta/ (recursively):

- meta/episodes/chunk-000/file-000.parquet
- meta/info.json
- meta/stats.json
- meta/tasks.parquet

```

[7]:
      relative_path  size_bytes  \
0  meta/episodes/chunk-000/file-000.parquet    72560
1                meta/info.json         3401
2                meta/stats.json         5914
3                meta/tasks.parquet         2246
4          data/chunk-000/file-000.parquet    369943

```

```

                                sha256
0  191998bffd2680c477a4000270cf943bc423ba5fb54ee8...
1  254909942a6cbfb4692a239e4d0aa5c68ec8eb16f81b65...
2  4ae86bed785e0f98914812e87736e216139a22b43cf2b9...
3  1040cdef3328ec4376587152647df8e725e80a04210c42...
4  579ad57e2454359fa9f2c0e83525991bd2e0305b7b2914...

```

Verification: data_card/file_hashes.csv was opened and confirmed to contain five real 64-character hex SHA-256 values (not placeholders). PINNED_REVISION (f641879e22172be7e8161d5e6c1503c2d2feb657) and the full hash table were copied into data_card/SOURCE_LICENSE_HASH.md.

Hash scope: file_hashes.csv covers the five files used in this audit (the four files under

`meta/` and the single `data/` parquet file). The pinned snapshot contains nine files; the other four (`.gitattributes`, the dataset card `README.md`, and the two video files) were not hashed here because they are not used in the analysis.

1.5 5. Loading all episodes into one tidy DataFrame

In this snapshot (LeRobot format v3.0), one parquet file under `data/` may pack several episodes together, so the number of files is not expected to equal the number of episodes. We concatenate all data files into a single frame-level table and separately load the episode-metadata table (`meta/episodes/`), which records each episode's global index range and per-episode statistics.

```
[8]: info_json = json.loads((local_dir / "meta" / "info.json").read_text())
    print(json.dumps(info_json, indent=2)[:1500])

{
  "codebase_version": "v3.0",
  "robot_type": "so100_follower",
  "total_episodes": 50,
  "total_frames": 11939,
  "total_tasks": 1,
  "chunks_size": 1000,
  "fps": 30,
  "splits": {
    "train": "0:50"
  },
  "data_path": "data/chunk-{chunk_index:03d}/file-{file_index:03d}.parquet",
  "video_path":
"videos/{video_key}/chunk-{chunk_index:03d}/file-{file_index:03d}.mp4",
  "features": {
    "action": {
      "dtype": "float32",
      "shape": [
        6
      ],
      "names": [
        "shoulder_pan.pos",
        "shoulder_lift.pos",
        "elbow_flex.pos",
        "wrist_flex.pos",
        "wrist_roll.pos",
        "gripper.pos"
      ],
      "fps": 30.0
    },
    "observation.state": {
      "dtype": "float32",
      "shape": [
        6
      ]
    }
  }
}
```

```

],
"names": [
    "shoulder_pan.pos",
    "shoulder_lift.pos",
    "elbow_flex.pos",
    "wrist_flex.pos",
    "wrist_roll.pos",
    "gripper.pos"
],
"fps": 30.0
},
"observation.images.up": {
    "dtype": "video",
    "shape": [
        480,
        640,
        3
    ],
    "names": [
        "height",
        "width",
        "channels"
    ],
    "info": {
        "video.height": 480,
        "video.width": 640,
        "video.codec": "av1",
        "video.pix_fmt": "yuv420p",
        "video.is_depth_map": false,
        "video.fps": 30,
        "video.channels": 3,
        "has_audio": false
    }
},
"observation.images.side": {
    "dtype": "video",

```

```

[9]: episodes_jsonl = local_dir / "meta" / "episodes.jsonl"
    episodes_dir = local_dir / "meta" / "episodes"

    if episodes_jsonl.is_file():
        print("Detected legacy format: meta/episodes.jsonl (single file)")
        episodes_meta = pd.read_json(episodes_jsonl, lines=True)
    elif episodes_dir.is_dir():
        print("Detected v3 format: meta/episodes/ (directory of parquet shards)")
        ep_parquet_files = sorted(episodes_dir.rglob("*.parquet"))

```

```

print(f" found {len(ep_parquet_files)} parquet shard(s)")
episodes_meta = pd.concat(
    [pd.read_parquet(p) for p in ep_parquet_files], ignore_index=True
)
else:
    raise FileNotFoundError("Neither meta/episodes.jsonl nor meta/episodes/_
↳found.")

print(episodes_meta.shape)
print(episodes_meta.columns.tolist())
episodes_meta.head()

```

Detected v3 format: meta/episodes/ (directory of parquet shards)

found 1 parquet shard(s)

(50, 62)

```

['episode_index', 'data/chunk_index', 'data/file_index', 'dataset_from_index',
'dataset_to_index', 'videos/observation.images.side/chunk_index',
'videos/observation.images.side/file_index',
'videos/observation.images.side/from_timestamp',
'videos/observation.images.side/to_timestamp',
'videos/observation.images.up/chunk_index',
'videos/observation.images.up/file_index',
'videos/observation.images.up/from_timestamp',
'videos/observation.images.up/to_timestamp', 'tasks', 'length',
'stats/action/min', 'stats/action/max', 'stats/action/mean', 'stats/action/std',
'stats/action/count', 'stats/observation.state/min',
'stats/observation.state/max', 'stats/observation.state/mean',
'stats/observation.state/std', 'stats/observation.state/count',
'stats/observation.images.up/min', 'stats/observation.images.up/max',
'stats/observation.images.up/mean', 'stats/observation.images.up/std',
'stats/observation.images.up/count', 'stats/observation.images.side/min',
'stats/observation.images.side/max', 'stats/observation.images.side/mean',
'stats/observation.images.side/std', 'stats/observation.images.side/count',
'stats/timestamp/min', 'stats/timestamp/max', 'stats/timestamp/mean',
'stats/timestamp/std', 'stats/timestamp/count', 'stats/frame_index/min',
'stats/frame_index/max', 'stats/frame_index/mean', 'stats/frame_index/std',
'stats/frame_index/count', 'stats/episode_index/min', 'stats/episode_index/max',
'stats/episode_index/mean', 'stats/episode_index/std',
'stats/episode_index/count', 'stats/index/min', 'stats/index/max',
'stats/index/mean', 'stats/index/std', 'stats/index/count',
'stats/task_index/min', 'stats/task_index/max', 'stats/task_index/mean',
'stats/task_index/std', 'stats/task_index/count', 'meta/episodes/chunk_index',
'meta/episodes/file_index']

```

```

[9]: episode_index  data/chunk_index  data/file_index  dataset_from_index  \
0                0                0                0                0
1                1                0                0                303

```

2	2	0	0	569
3	3	0	0	799
4	4	0	0	982

	dataset_to_index	videos/observation.images.side/chunk_index	\
0	303	0	
1	569	0	
2	799	0	
3	982	0	
4	1197	0	

	videos/observation.images.side/file_index	\
0	0	
1	0	
2	0	
3	0	
4	0	

	videos/observation.images.side/from_timestamp	\
0	0.000000	
1	10.100000	
2	18.966667	
3	26.633333	
4	32.733333	

	videos/observation.images.side/to_timestamp	\
0	10.100000	
1	18.966667	
2	26.633333	
3	32.733333	
4	39.900000	

	videos/observation.images.up/chunk_index	...	stats/index/mean	\
0	0	...	[151.0]	
1	0	...	[435.5]	
2	0	...	[683.5]	
3	0	...	[890.0]	
4	0	...	[1089.0]	

	stats/index/std	stats/index/count	stats/task_index/min	\
0	[87.46808941932291]	[303]	[0]	
1	[76.78704317786953]	[266]	[0]	
2	[66.39465339920075]	[230]	[0]	
3	[52.82676089508675]	[183]	[0]	
4	[62.0644825967316]	[215]	[0]	

	stats/task_index/max	stats/task_index/mean	stats/task_index/std	\
--	----------------------	-----------------------	----------------------	---

0	[0]	[0.0]	[0.0]
1	[0]	[0.0]	[0.0]
2	[0]	[0.0]	[0.0]
3	[0]	[0.0]	[0.0]
4	[0]	[0.0]	[0.0]

	stats/task_index/count	meta/episodes/chunk_index	meta/episodes/file_index
0	[303]	0	0
1	[266]	0	0
2	[230]	0	0
3	[183]	0	0
4	[215]	0	0

[5 rows x 62 columns]

```
[10]: parquet_files = sorted((local_dir / "data").rglob("*.parquet"))
print(f"Found {len(parquet_files)} parquet file(s) under data/ "
      f"(note: in this format one file may pack multiple episodes together, "
      f"so this count is NOT expected to equal total_episodes - that's fine).")

frames = [pq.read_table(fp).to_pandas() for fp in parquet_files]
df = pd.concat(frames, ignore_index=True)

n_episodes_seen = df["episode_index"].nunique()
print("Total frames loaded   :", len(df), " | declared total_frames:",
      info_json["total_frames"])
print("Distinct episodes seen:", n_episodes_seen, " | declared total_episodes:
      ", info_json["total_episodes"])
df.head()
```

Found 1 parquet file(s) under data/ (note: in this format one file may pack multiple episodes together, so this count is NOT expected to equal total_episodes - that's fine).

Total frames loaded : 11939 | declared total_frames: 11939

Distinct episodes seen: 50 | declared total_episodes: 50

```
[10]:                                     action \
0 [1.9117647, -99.41053, 99.54483, 74.84063, -48...
1 [1.9117647, -99.41053, 99.54483, 74.84063, -48...
2 [1.9117647, -99.41053, 99.54483, 74.84063, -48...
3 [1.9117647, -99.41053, 99.54483, 74.84063, -48...
4 [1.9117647, -99.41053, 99.54483, 74.84063, -48...

                                observation.state  timestamp  frame_index \
0 [1.9560878, -98.74372, 98.92425, 74.81984, -51...  0.000000          0
1 [1.9560878, -98.74372, 98.92425, 74.81984, -51...  0.033333          1
2 [1.9560878, -98.74372, 98.92425, 74.81984, -51...  0.066667          2
```

3	[1.9560878, -98.74372, 98.92425, 74.81984, -50...	0.100000	3
4	[1.9560878, -98.74372, 98.92425, 74.81984, -49...	0.133333	4

	episode_index	index	task_index
0	0	0	0
1	0	1	0
2	0	2	0
3	0	3	0
4	0	4	0

```
[11]: expected = {"robot_type": "so100_follower", "fps": 30, "total_episodes": 50,
                  "total_frames": 11939, "total_tasks": 1, "codebase_version": "v3.0"}
for k, v in expected.items():
    actual = info_json.get(k)
    print(f"{k:18s} expected={v!r:18} actual={actual!r:18} {'OK' if actual == v_
    ↪ else 'MISMATCH'}")
print("len(df) =", len(df), "| distinct episodes =", df["episode_index"].
    ↪ nunique())
print("task_index values:", sorted(df["task_index"].unique()))
print("action names:", info_json["features"]["action"]["names"])
```

robot_type	expected='so100_follower'	actual='so100_follower'	OK
fps	expected=30	actual=30	OK
total_episodes	expected=50	actual=50	OK
total_frames	expected=11939	actual=11939	OK
total_tasks	expected=1	actual=1	OK
codebase_version	expected='v3.0'	actual='v3.0'	OK

len(df) = 11939 | distinct episodes = 50
task_index values: [np.int64(0)]
action names: ['shoulder_pan.pos', 'shoulder_lift.pos', 'elbow_flex.pos',
'wrist_flex.pos', 'wrist_roll.pos', 'gripper.pos']

```
[12]: tasks_df = pd.read_parquet(local_dir / "meta" / "tasks.parquet")
print(tasks_df)
```

	task_index
pink lego brick into the transparent box	0

1.6 6. Schema audit

Cross-check the *declared* schema in meta/info.json (features block: dtype and shape per column) against what pandas/pyarrow actually loaded.

```
[13]: schema_report = []
for col, spec in info_json["features"].items():
    if col not in df.columns:
        schema_report.append({"column": col, "declared_dtype": spec["dtype"],
                              "declared_shape": spec["shape"],
```

```

        "status": "MISSING FROM PARQUET (likely a
↪video-only column)"}))
        continue
    sample_val = df[col].iloc[0]
    actual_len = len(sample_val) if hasattr(sample_val, "__len__") else 1
    declared_len = spec["shape"][0] if spec["shape"] else 1
    ok = (actual_len == declared_len)
    schema_report.append({
        "column": col,
        "declared_dtype": spec["dtype"],
        "declared_shape": spec["shape"],
        "actual_pandas_dtype": str(df[col].dtype),
        "actual_vector_len": actual_len,
        "shape_matches": ok,
    })

schema_df = pd.DataFrame(schema_report)
schema_df

```

```

[13]:
      column declared_dtype declared_shape actual_pandas_dtype \
0          action      float32           [6]           object
1  observation.state      float32           [6]           object
2  observation.images.up      video  [480, 640, 3]           NaN
3  observation.images.side      video  [480, 640, 3]           NaN
4          timestamp      float32           [1]      float32
5          frame_index      int64           [1]           int64
6          episode_index      int64           [1]           int64
7              index      int64           [1]           int64
8          task_index      int64           [1]           int64

      actual_vector_len shape_matches \
0              6.0           True
1              6.0           True
2              NaN           NaN
3              NaN           NaN
4              1.0           True
5              1.0           True
6              1.0           True
7              1.0           True
8              1.0           True

      status
0          NaN
1          NaN
2  MISSING FROM PARQUET (likely a video-only column)
3  MISSING FROM PARQUET (likely a video-only column)
4          NaN

```

5	NaN
6	NaN
7	NaN
8	NaN

Schema audit result: all seven tabular columns (`action`, `observation.state`, `timestamp`, `frame_index`, `episode_index`, `index`, `task_index`) matched their `meta/info.json`-declared shape in the first row of each column (`shape_matches = True` for all seven; only the first row was inspected, not every row). The two video columns (`observation.images.up`, `observation.images.side`) are, as expected, absent from the parquet tables — they exist only as separate (multi-episode) `.mp4` files and are explicitly out of scope for this audit (see Data Dictionary). Note that pandas reports `action/observation.state` as `dtype=object` rather than `float32`; this is expected and not a mismatch — each cell holds a length-6 array, so pandas can only type the *column* as `object`; the element dtype of those arrays was not separately verified here. No undeclared or unexpected columns were found in the parquet files.

1.7 7. Range and plausibility audit

We expand the 6-dimensional `action` and `observation.state` vectors into individual columns and inspect their empirical ranges. The physical unit is not documented in `meta/info.json`; based on the observed ranges (`shoulder_pan` about -93 to 88, `shoulder_lift` about -100 to 9, `elbow_flex` about 13 to 100, `wrist_flex` about 34 to 99, `wrist_roll` about -93 to -20, `gripper` about 0 to 33), this looks like a normalized position scale rather than raw degrees, but this is not confirmed by the dataset’s own documentation. There is no explicit `min/max` contract in `info.json`, so “plausible” here means: finite, not `NaN/inf`, and consistent between `action` and `observation.state` (which should track each other closely for a low-latency follower arm).

```
[14]: JOINT_NAMES = info_json["features"]["action"]["names"]

action_arr = np.stack(df["action"].to_numpy())
state_arr = np.stack(df["observation.state"].to_numpy())

action_wide = pd.DataFrame(action_arr, columns=[f"action.{n}" for n in JOINT_NAMES])
state_wide = pd.DataFrame(state_arr, columns=[f"state.{n}" for n in JOINT_NAMES])


range_report = pd.DataFrame({
    "min": pd.concat([action_wide, state_wide], axis=1).min(),
    "max": pd.concat([action_wide, state_wide], axis=1).max(),
    "mean": pd.concat([action_wide, state_wide], axis=1).mean(),
    "std": pd.concat([action_wide, state_wide], axis=1).std(),
    "n_nonfinite": (~np.isfinite(pd.concat([action_wide, state_wide], axis=1))).sum(),
})
range_report
```

```
[14]:
```

	min	max	mean	std \
action.shoulder_pan.pos	-93.455879	88.014709	8.021068	44.564873
action.shoulder_lift.pos	-100.000000	8.126316	-55.962349	36.486614
action.elbow_flex.pos	12.972235	100.000000	65.255653	29.013208
action.wrist_flex.pos	33.531662	99.490013	69.181854	13.238659
action.wrist_roll.pos	-92.771675	-20.000000	-53.419933	17.765167
action.gripper.pos	0.000000	32.998455	6.848944	8.999420
state.shoulder_pan.pos	-92.654694	88.103790	7.987031	44.493618
state.shoulder_lift.pos	-98.911224	8.793970	-55.180790	36.828537
state.elbow_flex.pos	16.808605	99.372482	66.776299	27.709524
state.wrist_flex.pos	34.039848	98.643494	69.252495	13.013702
state.wrist_roll.pos	-92.283272	-20.195360	-53.438889	17.703703
state.gripper.pos	1.208054	32.684563	8.112266	8.405106

	n_nonfinite
action.shoulder_pan.pos	0
action.shoulder_lift.pos	0
action.elbow_flex.pos	0
action.wrist_flex.pos	0
action.wrist_roll.pos	0
action.gripper.pos	0
state.shoulder_pan.pos	0
state.shoulder_lift.pos	0
state.elbow_flex.pos	0
state.wrist_flex.pos	0
state.wrist_roll.pos	0
state.gripper.pos	0

```
[15]: fig, axes = plt.subplots(2, 3, figsize=(12, 6), sharex=False)
for ax, name in zip(axes.ravel(), JOINT_NAMES):
    ax.hist(action_wide[f"action.{name}"], bins=60, alpha=0.6, label="action")
    ax.hist(state_wide[f"state.{name}"], bins=60, alpha=0.6, label="state")
    ax.set_title(name, fontsize=9)
    ax.legend(fontsize=7)
plt.suptitle("Distribution of each joint channel: commanded action vs measured_
↪state")
plt.tight_layout()
fig.savefig("../report/range_histograms.png", dpi=150)
plt.show()
```



Range audit result: no non-finite values were found in any of the 12 columns (`n_nonfinite` = 0 throughout). `action` and `observation.state` track each other closely for most joints — mean differences are small for `shoulder_pan.pos` (0.034), `wrist_roll.pos` (0.019) and `wrist_flex.pos` (0.071) — but diverge more for `elbow_flex.pos` (1.52), `gripper.pos` (1.26) and `shoulder_lift.pos` (0.78). For these three joints the mean of `action` - `state` is negative (-1.52, -1.26, -0.78), i.e. commanded values sit below measured values on average. A signed mean difference cannot distinguish tracking lag from a constant offset (calibration, steady-state error, or gripper contact when closing on an object); no lag or cross-correlation analysis was done.

Four boundary values were found in `action` (none in `observation.state`): `action.shoulder_lift.pos` reaches an exact minimum of -100.000000, `action.elbow_flex.pos` an exact maximum of 100.000000, `action.wrist_roll.pos` an exact maximum of -20.000000, and `action.gripper.pos` an exact minimum of 0.000000. The measured `observation.state` never reaches these exact values (state `shoulder_lift` min -98.91, `elbow_flex` max 99.37, `wrist_roll` max -20.20, `gripper` min 1.21). Landing exactly on round numbers in the command channel is consistent with the commanded action being clipped at configured limits; the cause (normalized-scale limit, calibration limit, or operator hardware) was not verified against LeRobot’s documentation or source. Anyone using `action` as a continuous regression target near these bounds should treat those values as saturated.

The histograms are visibly multi-modal for every joint, with tall, narrow spikes at `shoulder_lift` near -100 and `elbow_flex` near 100 (roughly 4,000 frames each in the plot), `gripper` near 1-2 (roughly 5,800 frames) and `wrist_flex` near 73 (roughly 2,000 frames). These spikes sit at or near the episode-start pose (Section 12), so a large share of frames is spent near that rest pose. The static-start repeated frames quantified in Section 9 are only about 3% of all frames (mean leading run 7.18 frames x 50 episodes, about 359 frames), far fewer than these spike heights, so the spikes must also include near-rest frames elsewhere in the episode (for example after returning to rest); that was not separately measured. In addition, `shoulder_pan` shows several separated peaks (around -88, -55, -22, 0, 45 and 85) that lie far from its start-pose values (-4 to 15). This multi-modality is therefore not explained by the rest pose and may reflect different target or placement

positions across episodes. This was not verified (the videos were not decoded), and it bears on the “fixed setup” assumption in Section 13.

The `gripper.pos` channel occupies a narrower numeric range (roughly 0-33) than `shoulder_pan` (span about 181), `shoulder_lift` (about 108) and the other joints (spans about 66-87). This is consistent with a different mechanism or scale for the gripper, but the units of all six channels are undocumented in `info.json`.

1.8 8. Missingness audit

Checked at two levels: (a) null/NaN values in the tabular columns actually present in the parquet files, and (b) structural “missingness” of a construct the dataset does **not** provide at all (a success/outcome label) — which is not a NaN in a column, but a column that was never collected.

```
[16]: tabular_cols = ["timestamp", "frame_index", "episode_index", "index",  
    ↪ "task_index"]  
na_counts = df[tabular_cols].isna().sum()  
na_counts_vecs = pd.Series({  
    "action": np.isnan(action_arr).sum(),  
    "observation.state": np.isnan(state_arr).sum(),  
})  
print("NaNs in scalar columns:\n", na_counts)  
print("\nNaNs in vector columns (summed over all 6 dims):\n", na_counts_vecs)  
  
has_success_label = any("success" in c.lower() or "reward" in c.lower() or  
    ↪ "done" in c.lower()  
                        for c in list(df.columns) + list(info_json["features"]  
    ↪ keys()))  
print("\nDoes the schema include any success/reward/done column?",  
    ↪ has_success_label)
```

NaNs in scalar columns:

timestamp	0
frame_index	0
episode_index	0
index	0
task_index	0
dtype:	int64

NaNs in vector columns (summed over all 6 dims):

action	0
observation.state	0
dtype:	int64

Does the schema include any success/reward/done column? False

Missingness audit result: zero NaN values were found across all scalar columns (`timestamp`, `frame_index`, `episode_index`, `index`, `task_index`) and both vector columns (`action`, `observation.state`) — 0 missing cells out of 11,939 frames × 17 numeric values per frame (5 scalar

columns + 6 action + 6 state channels).

Structural missingness finding: despite zero cell-level missingness, the dataset provides **no outcome label** at all — no **success**, **reward**, or **done-at-goal** flag (`has_success_label = False`). Every frame only records *what the operator/arm did*, never *whether the task actually succeeded*. This is a measurement-construct gap, not a data-quality bug: any claim about “task performance” made from this dataset alone would require an external annotation step this dataset does not include.

1.9 9. Duplicate-frame audit

At the start of many teleoperated recordings the arm is briefly stationary before the operator begins moving it. We check, per episode, how many **consecutive** frames have an exactly identical **action** vector to the previous frame (a proxy for “held / static” frames), and where in the episode they concentrate.

```
[17]: df_sorted = df.sort_values(["episode_index", "frame_index"]).
      ↪reset_index(drop=True)
      action_list = df_sorted["action"].tolist()

      is_dup = [False]
      for i in range(1, len(action_list)):
          same_episode = df_sorted.loc[i, "episode_index"] == df_sorted.loc[i - 1,
          ↪"episode_index"]
          same_action = same_episode and np.array_equal(action_list[i], action_list[i
          ↪- 1])
          is_dup.append(bool(same_action))

      df_sorted["is_exact_dup_of_prev"] = is_dup
      dup_rate_overall = df_sorted["is_exact_dup_of_prev"].mean()
      print(f"Overall fraction of frames that are an exact duplicate of the previous
      ↪frame: {dup_rate_overall:.2%}")

      dup_by_episode = df_sorted.groupby("episode_index")["is_exact_dup_of_prev"].
      ↪mean()
      print("\nPer-episode duplicate-frame rate (head):")
      print(dup_by_episode.head(10))
```

Overall fraction of frames that are an exact duplicate of the previous frame:
14.13%

Per-episode duplicate-frame rate (head):

episode_index	
0	0.323432
1	0.120301
2	0.130435
3	0.142077
4	0.144186
5	0.173160
6	0.174107

```

7    0.157635
8    0.234513
9    0.151961
Name: is_exact_dup_of_prev, dtype: float64

```

```

[18]: # Where in the episode do duplicates concentrate? (expect: near frame_index 0)
def leading_run_length(is_dup_arr):
    # is_dup_arr[0] is trivially False (no previous frame within this episode),
    # so the *leading run of duplicate frames* starts being meaningful from
    ↪ index 1:
    # count consecutive True values starting right after the (always-False)
    ↪ first frame.
    count = 0
    for v in is_dup_arr[1:]:
        if v:
            count += 1
        else:
            break
    return count

leading_dup_run = (
    df_sorted.groupby("episode_index")["is_exact_dup_of_prev"]
    .apply(lambda s: leading_run_length(s.to_numpy()))
)

print("Length of the *leading* run of duplicate frames per episode (head):")
print(leading_dup_run.head(10))
print("\nMedian leading run:", leading_dup_run.median(), "frames")
print("Mean leading run  :", leading_dup_run.mean(), "frames")
print("Episodes with a leading run > 0:", (leading_dup_run > 0).sum(), "/",
    ↪ len(leading_dup_run))
print("Length of the *leading* run of duplicate frames per episode (head):")
print(leading_dup_run.head(10).astype(int))

fig, ax = plt.subplots(figsize=(7, 4))
ax.hist(leading_dup_run, bins=20)
ax.set_xlabel("Leading duplicate-frame run length (frames held static before
    ↪ first move)")
ax.set_ylabel("Number of episodes")
ax.set_title("Static-start duplicate frames across the 50 episodes")
plt.tight_layout()
fig.savefig("../report/duplicate_frames.png", dpi=150)
plt.show()

```

```

Length of the *leading* run of duplicate frames per episode (head):
episode_index
0    76
1     5
2     0

```

3	6
4	6
5	10
6	8
7	12
8	19
9	1

Name: is_exact_dup_of_prev, dtype: int64

Median leading run: 3.0 frames

Mean leading run : 7.18 frames


Episodes with a leading run > 0: 37 / 50

Length of the *leading* run of duplicate frames per episode (head):

episode_index

0	76
1	5
2	0
3	6
4	6
5	10
6	8
7	12
8	19
9	1

Name: is_exact_dup_of_prev, dtype: int64



```
[19]: print(dup_by_episode.describe())
```

```
count      50.000000
mean       0.140174
std        0.047673
min        0.059113
25%        0.112323
50%        0.130435
75%        0.153662
max        0.323432
Name: is_exact_dup_of_prev, dtype: float64
```

Duplicate-frame audit result: overall, 14.13% of all 11,939 frames have an `action` vector identical to the previous frame’s within the same episode (“repeated-action” frames). `observation.state` and the video frames of these rows were not compared, so they are not necessarily duplicates in the full-row sense. Per-episode, this rate ranges from 5.91% to 32.34% (mean 14.02%, median 13.04%, std 4.77%) — **every single episode has a non-trivial duplicate rate, with a floor of ~6%**, not just the few episodes with unusually long static starts.

Considering only the *leading* run of duplicates at the start of each episode (before the operator’s first move), the median run length is 3 frames (0.1 s at 30 fps) and the mean is 7.18 frames (0.24 s), with 37 of the 50 episodes (74%) showing at least one leading duplicate frame; one clear outlier, episode 0, has a 76-frame (2.53 s) static start. But a mean leading run of 7.18 frames across a ~239-frame average episode accounts for only ~3% of frames — far short of the observed ~14% overall duplicate rate. This gap shows that **most duplicate frames are not confined to the episode start**: they also occur later in the episode; mid-episode pauses and trailing static frames at the end of an episode are both plausible. Only the leading run was measured, so where the remaining ~11 percentage points fall was not established.

Implication: if a frame-level train/test split is used, both splits will contain a non-trivial share (~14%, with a per-episode floor of ~6%) of near-identical “held” frames. Whether these repeated-action frames also have near-identical `observation.state` (which would make test rows easy to predict from training rows) was not measured here; that is part of the leakage question in Section 11. Roughly one in seven frames may be an easy, low-variance prediction target, which could make an evaluation metric look better than it would on the dynamic portion of the trajectory. This should be disclosed alongside the leakage result rather than conflated with it.

1.10 10. Anomaly audit: cross-source consistency

The dataset states each episode’s extent in two independent places: the episode-metadata table (`meta/episodes/`, from which we derive length as `dataset_to_index - dataset_from_index`) and the frame-level parquet files (row count per `episode_index`). They should agree exactly. We also check that `timestamp` increases at a constant ~1/30 s step within each episode (a proxy for dropped/duplicated frames or clock jitter).

```
[20]: # Derive each episode's length from its global index range (dataset_to_index is
# exclusive, matching Python slicing) as an independent cross-check against the
      ↪parquet
```

```

# row counts. (The v3 episode metadata also contains a `length` column, not
↳used here.)
ep = episodes_meta.set_index("episode_index")
declared_lengths = (ep["dataset_to_index"] - ep["dataset_from_index"]).
↳rename("declared_length")
actual_lengths = df_sorted.groupby("episode_index").size().
↳rename("actual_length")

length_check = pd.concat([declared_lengths, actual_lengths], axis=1)
length_check["match"] = length_check["declared_length"] ==
↳length_check["actual_length"]
print("Episodes with a length mismatch between meta/episodes (derived) and
↳parquet:")
print(length_check[~length_check["match"]])
print(f"\n{length_check['match'].sum()} / {len(length_check)} episodes match
↳exactly.")

```

Episodes with a length mismatch between meta/episodes (derived) and parquet:
Empty DataFrame
Columns: [declared_length, actual_length, match]
Index: []

50 / 50 episodes match exactly.

```

[21]: FPS = info_json["fps"]
expected_dt = 1.0 / FPS
tolerance = expected_dt * 0.05 # 5% tolerance

anomalies = []
for ep, g in df_sorted.groupby("episode_index"):
    ts = g.sort_values("frame_index")["timestamp"].to_numpy()
    dt = np.diff(ts)
    bad = np.abs(dt - expected_dt) > tolerance
    if bad.any():
        anomalies.append({"episode_index": ep, "n_irregular_steps": int(bad.
↳sum()),
                                "max_abs_deviation_s": float(np.max(np.abs(dt -
↳expected_dt)))})

anomaly_df = pd.DataFrame(anomalies)
print(f"Episodes with at least one irregular inter-frame timestamp step "
    f"(tolerance = {tolerance*1000:.1f} ms): {len(anomaly_df)} /
↳{episodes_meta.shape[0]}")
anomaly_df.head(10)

```

Episodes with at least one irregular inter-frame timestamp step (tolerance = 1.7 ms): 0 / 50

```
[21]: Empty DataFrame
      Columns: []
      Index: []
```

Anomaly audit result: both cross-source consistency checks passed with zero discrepancies. All 50 episodes' declared frame counts (derived from `dataset_to_index - dataset_from_index` in the v3 episode metadata) matched the actual parquet row counts exactly (50/50, 0 mismatches) — no evidence of a mismatch between the episode metadata and the parquet packaging. Timestamp spacing also stayed within the 5% tolerance (1.7 ms) of the declared 1/30 s step for every single episode (0/50 episodes with an irregular step). The `timestamp` column may be a nominal value (`frame_index / fps`) rather than a measured capture time; if so, this check confirms internal consistency only and cannot detect real clock jitter or dropped frames. This was not verified. No anomalies were found under either check; this is reported as a negative result. Both checks test consistency between metadata and parquet only; value-level outlier detection (for example frame-to-frame jumps in `action` or `observation.state`) was not performed.

1.11 11. Leakage check: frame-level vs episode-level split

We fit a deliberately simple, non-tunable baseline — a k-nearest-neighbors regressor — to predict `action` from `observation.state`, under two splitting protocols, and compare held-out MSE. k-NN is chosen because it is simple, has no training-time tuning, and is expected to be particularly sensitive to near-duplicate neighbors. It is therefore a demonstration of a sensitive case, not a typical model, and the size of the gap should not be assumed to carry over to other models.

```
[22]: X = state_arr # (n_frames, 6)
      y = action_arr # (n_frames, 6)
      groups = df_sorted["episode_index"].to_numpy()

      def evaluate_split(train_idx, test_idx, k=1):
          model = KNeighborsRegressor(n_neighbors=k)
          model.fit(X[train_idx], y[train_idx])
          pred = model.predict(X[test_idx])
          return mean_squared_error(y[test_idx], pred)

      results = []
      for k in (1, 3, 5):
          # (a) frame-level random split -- ignores episode grouping
          frame_splitter = ShuffleSplit(n_splits=5, test_size=0.2, random_state=SEED)
          frame_mses = [evaluate_split(tr, te, k=k) for tr, te in frame_splitter.
                        ↪split(X)]

          # (b) episode-level grouped split -- an episode never appears in both sides
          group_splitter = GroupShuffleSplit(n_splits=5, test_size=0.2,
          ↪random_state=SEED)
          group_mses = [evaluate_split(tr, te, k=k) for tr, te in group_splitter.
                        ↪split(X, y, groups)]

          results.append({
```

```

        "k": k,
        "frame_level_MSE_mean": np.mean(frame_mses),
        "frame_level_MSE_std": np.std(frame_mses),
        "episode_level_MSE_mean": np.mean(group_mses),
        "episode_level_MSE_std": np.std(group_mses),
        "leakage_gap (episode - frame)": np.mean(group_mses) - np.
↪mean(frame_mses),
    })

leakage_report = pd.DataFrame(results)
leakage_report

```

```

[22]:      k  frame_level_MSE_mean  frame_level_MSE_std  episode_level_MSE_mean \
0    1             5.353241             0.223655             18.529543
1    3             4.421131             0.138875             14.982736
2    5             5.268494             0.187832             13.661621

      episode_level_MSE_std  leakage_gap (episode - frame)
0              2.507665              13.176302
1              2.385420              10.561605
2              2.415241               8.393126

```

```

[23]: # Naive "no-op" baseline for context: predict the previous frame's action as ↪
↪this
# frame's action. This calibrates how autocorrelated the series is on its own,
# independent of any train/test split.
prev_action = np.roll(action_arr, shift=1, axis=0)
prev_action[0] = action_arr[0]
same_episode_mask = (df_sorted["frame_index"] > 0).to_numpy()
naive_mse = mean_squared_error(action_arr[same_episode_mask], ↪
↪prev_action[same_episode_mask])
print(f"Naive 'carry previous action forward' baseline MSE (context only): ↪
↪{naive_mse:.4f}")

```

Naive 'carry previous action forward' baseline MSE (context only): 1.8810

```


[24]: fig, ax = plt.subplots(figsize=(7, 4.5))
xk = leakage_report["k"]
ax.plot(xk, leakage_report["frame_level_MSE_mean"], "o-", label="Frame-level ↪
↪split (leaky)")
ax.plot(xk, leakage_report["episode_level_MSE_mean"], "o-", ↪
↪label="Episode-level split (grouped)")
ax.axhline(naive_mse, color="gray", linestyle="--", linewidth=1, label="Naive ↪
↪carry-forward baseline")
ax.set_xlabel("k (number of neighbors)")
ax.set_ylabel("Held-out MSE (action space)")
ax.set_title("Leakage effect: frame-level vs episode-level train/test split")

```

```

ax.legend(fontsize=8)
plt.tight_layout()
fig.savefig("../report/leakage_comparison.png", dpi=150)
plt.show()

```



Leakage audit result: at $k=1$, the frame-level (leaky) split yields a mean held-out MSE of 5.35, while the episode-level (grouped) split yields 18.53 — a leakage gap of 13.18, meaning the naive frame-level evaluation understates the true error by a factor of $3.46\times$. The same direction and a comparable magnitude hold at $k=3$ ($3.39\times$) and $k=5$ ($2.59\times$): episode-level error is consistently 2.6–3.5 $\times$ higher than frame-level error across all tested k . This is consistent with the hypothesis in Section 2: a random frame-level split puts frames from the same episode on both sides of the boundary, and its error estimate is substantially lower (more optimistic) than the episode-grouped estimate for this model. The mechanism (near-duplicate temporal neighbors crossing the split) was not tested directly — no buffered split or nearest-neighbor-distance analysis was run — so it is a plausible explanation, not an established one. The grouped-split MSE also varies noticeably across the five splits (std about 2.4–2.5), so the ratios are approximate. The result comes from one k -NN model on 50 episodes and does not by itself quantify how optimistic a frame-level benchmark would be for other models. Note also that the declared split in `info.json` (`train: 0:50`) simply contains all 50 episodes, i.e. the dataset defines no held-out set, so the evaluation protocol has to be chosen by the user.

A secondary, unexpected finding: the naive “carry the previous action forward” baseline (MSE = 1.88) actually outperforms the k -NN state $\rightarrow$ action model under **both** splitting protocols and all tested k (best k -NN result: 4.42 at $k=3$, frame-level). This does not undermine the leakage conclusion — the comparison there is between the same model under two splits, not between models

— but it is worth disclosing honestly: this comparison is not like-for-like. The carry-forward baseline uses the previous *action* (past label information) as its input, whereas k-NN sees only the current `observation.state`, and the baseline was computed over all frames rather than on the same test folds. The result shows how strongly *action* is autocorrelated and only suggests (does not prove) that a single state snapshot is a weaker predictor than temporal persistence. Baselines that use the same input as k-NN (for example predicting `action = observation.state`, or the training mean) were not run.

1.12 12. Coverage check: how similar are the 50 starting configurations?


A dataset with 50 episodes only supports claims about the specific slice of state-space it actually covers. We plot each episode's **initial** `observation.state` (`frame_index == 0`) projected onto its two most variable dimensions, as a quick, honest look at how much the arm's starting configuration actually varies across episodes.

```
[25]: start_states = df_sorted[df_sorted["frame_index"] == 0].
      ↪sort_values("episode_index")
start_arr = np.stack(start_states["observation.state"].to_numpy())

variances = start_arr.var(axis=0)
dim_order = np.argsort(variances)[::-1]
d0, d1 = dim_order[0], dim_order[1]

fig, ax = plt.subplots(figsize=(6, 5))
ax.scatter(start_arr[:, d0], start_arr[:, d1], s=40)
ax.set_xlabel(f"{JOINT_NAMES[d0]} (start)")
ax.set_ylabel(f"{JOINT_NAMES[d1]} (start)")
ax.set_title(f"Episode-start state coverage (n={len(start_arr)} episodes)")
plt.tight_layout()
fig.savefig("../report/start_state_coverage.png", dpi=150)
plt.show()

print("Per-dimension std of episode-start state:")
for name, v in zip(JOINT_NAMES, start_arr.std(axis=0)):
    print(f"  {name:20s}: std = {v:.3f}")
```



Per-dimension std of episode-start state:

```

shoulder_pan.pos    : std = 4.026
shoulder_lift.pos   : std = 0.267
elbow_flex.pos       : std = 0.136
wrist_flex.pos       : std = 0.977
wrist_roll.pos       : std = 4.411
gripper.pos          : std = 0.553

```

```

[26]: print("start min:", start_arr.min(axis=0).round(2))
      print("start max:", start_arr.max(axis=0).round(2))

```

```

start min: [ -4.03 -98.91  98.66  68.21 -63.27   1.21]
start max: [ 15.05 -97.4   99.28  74.82 -42.08   5.3 ]

```

Coverage audit result: episode-start states are tightly clustered, not spread across each joint's observed operating range. Four of the six joints (`elbow_flex.pos` std=0.136, `shoulder_lift.pos` std=0.267, `gripper.pos` std=0.553, `wrist_flex.pos` std=0.977) are essentially constant at episode start, versus a dataset-wide `observation.state` std of 27.7 (`elbow_flex`), 36.8 (`shoulder_lift`), 8.4 (`gripper`) and 13.0 (`wrist_flex`) (Section 7's range report) — the start-of-episode spread is roughly 13x (`wrist_flex`) to 204x (`elbow_flex`) smaller than the spread across the whole dataset. Only two

joints vary meaningfully at episode start: `shoulder_pan.pos` (std=4.026, observed range -4.03 to 15.05) and `wrist_roll.pos` (std=4.411, observed range -63.27 to -42.08), and even these span a modest band (about 19 and 21 units) relative to their observed dataset-wide ranges of about 181 and 72. The scatter plot shows the 50 starting configurations spread over a modest cloud, with two visibly separated episodes (one with `shoulder_pan` about 15, one with `wrist_roll` about -42).

This is consistent with, but does not by itself establish, the bias counterexample in Section 13: the arm starts from nearly the same pose in almost every episode, with only two of six joints showing any real initial-condition variation. A near-identical start pose is most likely a consequence of the recording protocol (a fixed rest pose) and does not show that objects, box position, lighting or operators were fixed; the multi-modal `shoulder_pan` distribution in Section 7 hints that target positions may vary. A model fit on this sample has little basis for generalizing to a meaningfully different starting configuration, because the sample essentially never demonstrates one.

1.13 13. Bias/measurement counterexample and a prohibited claim

Counterexample (concrete, falsifiable). This dataset records a *single task* (“pink lego brick into the transparent box”) performed by an *undisclosed number of teleoperators*, in a setup whose variation across episodes is undisclosed (the metadata shows one robot type, one pair of camera streams and one task string for all 50 episodes, `total_tasks = 1`, but says nothing about object, box, lighting or camera-position variation; the multi-modal `shoulder_pan` distribution in Section 7 hints that target positions may in fact vary). If the true population of interest is “SO-100/SO-101 pick-and-place behavior in general” (varied objects, containers, lighting, clutter, operators), this sample is a convenience sample of that population with unknown coverage of objects, containers, lighting, clutter and operators, not a random or stratified one. A model fit on it can be expected to fit this specific setup’s action distribution well while telling us little about performance on a different object, a different box position, or a different operator’s motion style — the sample is not exchangeable with that broader population.

Measurement counterexample. `action` is a commanded signal, not an outcome. In this dataset it is saturated at round-number values on four channels (Section 7) and repeats the previous frame’s value in 14.13% of frames (Section 9). A model that reproduces `action` well is reproducing the operator’s command signal (including its saturation and pauses), which is not evidence that the task was completed.

Prohibited claim (do not make this claim from this dataset alone): > “Because a policy trained on `lerobot/svla_so101_pickplace` achieves low held-out > prediction error, it will generalize to picking up new objects or operating in new > environments.” > “Because the frame-level held-out error on this dataset is low, a model will perform > comparably on a new episode.”

“Low action-prediction MSE means the robot completes the pick-and-place task.”

This is not supported: (1) there is no success/outcome label to even define “generalization” behaviorally (Section 8); (2) the sample covers one task, its episode-start poses are nearly identical (Section 12; this reflects the recording protocol and does not by itself show that objects or environments were fixed), and variation in objects, box position, lighting and operators is undisclosed; (3) low error under a frame-level split can be an optimistic estimate when frames from the same episode sit on both sides of the split (Section 11, for the k-NN baseline error was 2.6-3.5x higher under an episode-grouped split). The two added claims above fail for the same reasons: the first is contradicted by Section 11, the second by the absence of any success label (Section 8). Any claim

about out-of-distribution generalization would require new data collected under deliberately varied conditions, evaluated with an episode-level (or condition-level) held-out split.

1.14 14. Summary of findings

- **Schema:** all 7 tabular columns matched their `meta/info.json`-declared shape in the first row of each column; the 2 video columns are, as expected, absent from the parquet tables and were explicitly excluded from this audit’s scope.
- **Range:** no non-finite values anywhere (0/11,939 frames $\times$ 12 channels). `action` and `observation.state` track closely for most joints (absolute mean difference <0.1 for shoulder_pan/wrist_roll/wrist_flex) but differ more for elbow_flex (-1.52), gripper (-1.26) and shoulder_lift (-0.78). Four `action` channels sit exactly on round-number bounds (shoulder_lift min -100, elbow_flex max 100, wrist_roll max -20, gripper min 0) while `observation.state` never does, consistent with saturated commands; the cause was not verified.
- **Missingness:** 0 NaN cells overall, but a structural gap: no success/reward/done label exists anywhere in the schema — outcome cannot be assessed from this dataset alone.
- **Duplicates:** 14.13% of all frames are repeated-action frames (`action` identical to the previous frame) (per-episode range 5.9%–32.3%, every episode 5.9%). Only ~3% of frames are attributable to episode-start static holds (median 3, mean 7.18 leading frames per episode); the remaining ~11 percentage points occur outside the leading run (location not measured).
- **Anomalies:** 0 discrepancies — all 50 episodes’ declared frame counts matched parquet row counts exactly, and 0/50 episodes showed timestamp irregularities beyond a 5% (1.7 ms) tolerance on the declared 1/30 s step. (the timestamp check may be tautological if `timestamp` is nominal; not verified) Reported as a genuine negative result.
- **Leakage gap:** substantial and consistent in direction across $k=1,3,5$ — episode-level MSE is $2.6\times\text{--}3.5\times$ higher than frame-level MSE (13.18 absolute gap at $k=1$: 5.35 vs 18.53). For this k -NN baseline, the frame-level split reported error roughly 2.6-3.5x lower than the episode-grouped split. (Secondary finding: a naive carry-forward baseline, MSE=1.88, beats the k -NN model under every split/ k tested — disclosed for honesty, does not affect the leakage conclusion.)
- **Coverage / bias:** 4 of 6 joints are essentially constant at episode start (std 0.14–0.98, versus a dataset-wide `observation.state` std of 8.4–36.8 for the same four joints); only 2 joints vary at episode start, over about 19 (shoulder_pan) and 21 (wrist_roll) units, well under half of their observed dataset-wide ranges (about 181 and 72). Episodes start from a near-identical pose (likely the recording protocol’s rest pose); this alone does not show that objects or environments were fixed.

Answer to the testable question (Section 2): No — for this dataset a random frame-level split should not be used to evaluate an action-prediction baseline. For the k -NN baseline it produced error estimates $2.6\times\text{--}3.5\times$ lower than the episode-grouped split across $k=1,3,5$, which is consistent with optimistic bias from temporally adjacent, near-identical frames of the same episode appearing on both sides of the split (the mechanism itself was not tested separately). An episode-level (grouped) split is therefore the appropriate protocol for a trustworthy held-out evaluation. Separately, the dataset’s single-task, near-fixed-start sample (Sections 12–13) means that even a correctly evaluated model’s error says little about generalization beyond this specific setup.

1.15 15. Limitations and what this audit does not establish

- Pixel content of the two video streams was not decoded or audited (Section 6) — a full audit would additionally check for lighting drift, occlusion, and camera desync using the `.mp4` files.
- Provenance gaps (operator count, environment variability) were not resolved because the dataset card does not disclose them; they are reported as open questions, not assumed away.
- The k-NN leakage demonstration is a diagnostic, not a claim about the best possible imitation-learning model for this data.
- The leakage mechanism (near-duplicate neighbors) was not tested directly; only one model class (k-NN) and five overlapping random splits were used, with no confidence intervals, and no same-input trivial baselines were run.
- Range boundary values and the `timestamp` column were not verified against LeRobot’s documentation or source; units of all channels are undocumented.
- Episode 0 (76-frame static start, 32% repeated-action rate) was not excluded or tested for its effect on the results.
- Hashes cover five of the nine files in the pinned snapshot; the video files were neither hashed nor decoded.
- Results come from the runs recorded in this repository; the reproducibility evidence (second run, command-line hash check) is in `ai_use_log/AI_USE_LOG.md`.

1.16 16. AI use and independent verification

I used Claude (Anthropic) to plan the audit and write the notebook scaffold, and GPT 5.6 Luna as a second reviewer of the finished repository. Claude was also used to review the repository against the assignment brief and rubric (it produced issue lists only, no rewriting). Neither had network access to run this notebook against the real data — every number above comes from my own runs. Running the AI-written code against the actual pinned snapshot surfaced several real defects that I diagnosed and had fixed: an `IsADirectoryError` and two wrong format assumptions from the dataset’s actual LeRobot v3.0 layout (vs. the v2.1 layout the webpage suggested during planning), a `cumprod()` logic bug in the duplicate-run calculation (caught because its output silently contradicted the reported duplicate rate), a reproducibility-breaking `assert` in the revision-pinning code (found by GPT’s review), a mis-stated physical unit and two mis-read plot values in Sections 7 and 12 (caught by checking the drafted text against `info_json` and the code’s own printed min/max), a column-misaligned CSV in the data dictionary, and a cell placed before its dependency was defined (only surfaced by a full `Restart & Run All`, not by running cells individually). Two complete clean reruns produced identical `key_metrics.json` (to 6 decimal places) and identical file hashes; all five hashed files were also independently recomputed with command-line `shasum`. Full disclosure, including what AI-drafted text I kept largely as written and why: `ai_use_log/AI_USE_LOG.md`.

```
[27]: key_metrics = {
    "dup_rate_overall": round(float(dup_rate_overall), 6),
    "dup_by_episode_describe": dup_by_episode.describe().round(6).to_dict(),
    "leading_dup_run_median": float(leading_dup_run.median()),
    "leading_dup_run_mean": round(float(leading_dup_run.mean()), 6),
    "episodes_with_leading_run": int((leading_dup_run > 0).sum()),
    "length_match_count": int(length_check["match"].sum()),
    "timestamp_anomaly_episodes": int(len(anomaly_df)),
    "leakage_report": leakage_report.round(6).to_dict(orient="records"),
```

```
    "naive_mse": round(float(naive_mse), 6),  
    "start_state_std": [round(float(v), 6) for v in start_arr.std(axis=0)],  
}  
with open("../report/key_metrics.json", "w") as f:  
    json.dump(key_metrics, f, indent=2)  
print("Wrote ../report/key_metrics.json")
```

Wrote ../report/key_metrics.json